



UNIVERSITY OF
SAN FRANCISCO

CHANGE THE WORLD FROM HERE

Features & Visualisations

Machine Learning



What is Machine Learning?

- Use *features* (X) to describe a dataset of *instances*
 - Instances can be things, events, categories, etc.
 - Features have many names, including (predictor) variables, stimulus, etc.
- Unsupervised Learning
 - Principal Components Analysis = describe variance in dataset
 - Clustering = group / segment data instances
- Supervised Learning
 - One part of data is a target (Y)
 - We expect that Y is correlated to each X by some function
 - Tasks:
 - We need to choose our features (X) so they correlated with Y
 - We need to choose the function on each X that minimises the error in predicting Y



Machine learning process

- Goal
 - Unsupervised: describe X
 - Supervised: use X to predict Y
- Process
 - Determine the problem (classification?) and choose one or more performance metrics (accuracy?)
 - Collect data?
 - Visualise the data — identify types of data and plot them, clean outliers, etc.
 - Transform data / engineer features?
 - Implement and tune model(s); compare efficacy of each model using chosen metric(s)... or vs. baseline

ML in production: expectation

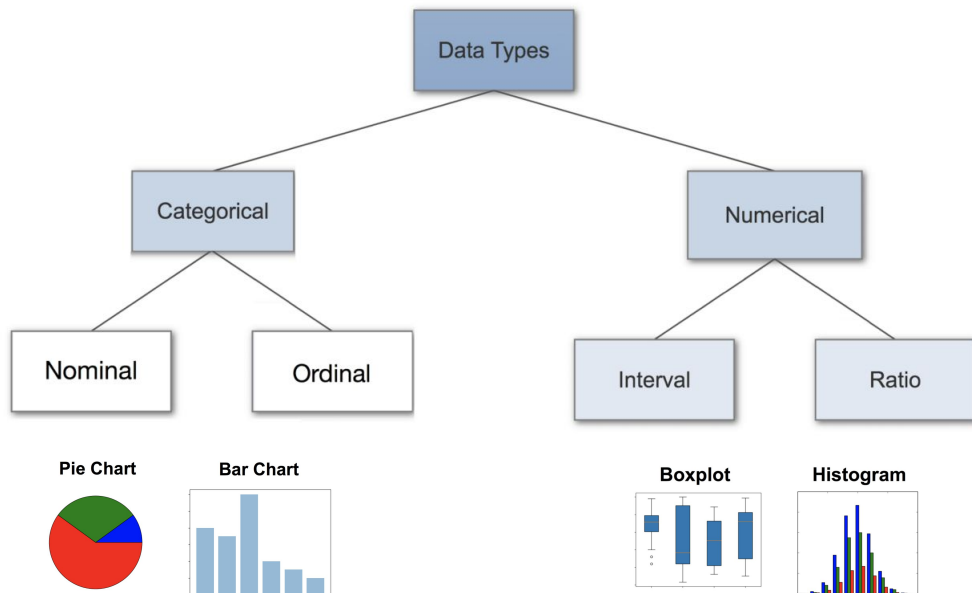
1. Collect data
2. Train model
3. Deploy model
- 4.



<https://twitter.com/chipro/status/1348265019012743169>



Types of data



Has / Can	Nominal	Ordinal	Interval	Ratio
Count / Frequency / Proportion	✓	✓	✓	✓
Mode, Median		✓	✓	✓
Order of values is known		✓	✓	✓
Mean, stdev of values			✓	✓
Add / subtract values			✓	✓
Multiply / divide values				✓
True zero				✓

<https://towardsdatascience.com/data-types-in-statistics-347e152e8bee>



Feature Engineering

- The art of creating good feature variables (X) from raw data
- Good feature set
 - Describes or represents the target well with fewest values
 - Leads to good models (better results)
 - Requires less complex algorithms
- Requires a lot of domain knowledge?
 - Listen to Subject Matter Experts (SMEs), knowing:
 - SMEs are often overestimate how predictive some data is
 - SMEs often miss some important predictors (or combinations of predictors)
 - SMEs often lack one or more data transformation steps



(One) Feature Engineering Process

- Get SME input
- Brainstorm features
 - Look at what other people have done*
 - Encourage wild ideas (especially working in teams) — and get a lot of ideas
- Decide which one(s) to use in the model
 - Judge effort vs. expected return in power of model
 - Go for novelty so that the model becomes more powerful
- Implement the features above
- Study the impact of the implemented features
- Repeat, repeat, repeat



Algorithms vs. outliers

Class	Algorithm	Sensitive to outliers?
Unsupervised	K-Means	Sensitive
	Hierarchical Clustering	Sensitive
	PCA	Sensitive
Regression	Linear Regression	Sensitive
Classification	Logistic Regression	Sensitive
	K-Nearest Neighbours	Not sensitive
	Naive Bayes, SVM	Not sensitive
	Decision Trees, Random Forest, Boosted Trees	Not sensitive
	Neural Networks	Sensitive



Data → Features (1)

- Imputation

- Numerical: default value (eg. 0), mean, median, etc.
- Categorical: most frequent value

- Outlier detection


- Standard deviation is often used for detecting outliers
- Common to project data to normal distribution and flagging values $\pm 3x$ std. Dev

- Binning

- Combining range of numerical values or multiple categorical values
- Use carefully since it loses (important?) information

- Mathematical transformation

- *Log transform* can make small differences more significant
- Sine / cosine to cyclical data (eg. dates) can show that Jan & Dec are closer than May & Jul



Canada	→ North America
Iran	→ Asia
Thailand	→ Asia
USA	→ North America



Data —> Features (2)

- One-hot encoding
 - Encoded categorical data may look related — eg. 1 = Alabama, 2 = Alaska
 - 1-hot encoding transforms data to multiple unrelated binary columns
- Splitting
 - Taking a string (eg. date: 2021-01-24) and extracting important parts
 - Also typical for NLP
- Lookup / External
 - Dates —> holidays, weekday, etc.
 - External sources are often useful
- Feature combinations
 - Add, subtract, multiply or divide two features to form a third
 - Remove features which are not useful




Why use visualisations?

- Difficult to see data patterns, trends, etc. in “wall of numbers” charts
- We can sometimes determine functions relating X to Y
- We can see outliers, possible errors in data, etc.

a

I		II		III		IV	
x	y	x	y	x	y	x	y
10	8.04	10	9.14	10	7.46	8	6.58
8	6.95	8	8.14	8	6.77	8	5.76
13	7.58	13	8.74	13	12.74	8	7.71
9	8.81	9	8.77	9	7.11	8	8.84
11	8.33	11	9.26	11	7.81	8	8.47
14	9.96	14	8.10	14	8.84	8	7.04
6	7.24	6	6.13	6	6.08	8	5.25
4	4.26	4	3.10	4	5.39	19	12.5
12	10.84	12	9.13	12	8.15	8	5.56
7	4.82	7	7.26	7	6.42	8	7.91
5	5.68	5	4.74	5	5.73	8	6.89

b



<https://towardsdatascience.com/9-data-visualization-tools-that-you-cannot-miss-in-2019-3ff23222a927>



What do visualisations need?

Title — Count of unique speakers outside the USA

*Labels for each
axis (or scale) —
12 speakers*






Tufte

- Known for...
 - Data visualisation pioneer
 - “Formally documented” the sparkline chart
- Principles for graphic design
 - Graphical Integrity
 - Chart axes must be labelled
 - Scale must be consistent
 - Minimalism
 - Use as little “ink” to show the data
 - Use few (no?) graphical effects — instead, show how the data varies

Example sparklines in small multiple



<https://en.wikipedia.org/wiki/Sparkline>



Why matplotlib / seaborn?

- matplotlib
 - Basic plotting library, default for python
 - Widely used
- seaborn
 - A layer on top of matplotlib
 - Provides additional functionality and improved aesthetics
- Not ggplot
 - A copy (port?) of package in R (ggplot2)
 - Not pythonic, therefore not intuitive for data people who are stronger in python
 - ... but preferred by many for its OO-treatment of charts



Set up matplotlib & seaborn

- Basic syntax

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
%matplotlib inline
```

- Common plot types:


- Bar plot (“countplot”)
- Box plot
- Line plot
- Scatterplot



Example

- Countplot (bar plot)
 - x-axis = category labels
 - y-axis is frequency count of category
- Basic syntax

```
sns.countplot(x = df['LABEL'], data = df)
```



- Example
 - 2018 [FARS](#) ACCIDENT.CSV
 - Best day of week to be on the road is Wednesday?



Some parameters to countplot

- Add additional dimensions (as “hue”)

```
sns.countplot(x = df['LABEL'], hue=df['WEATHER1'], data = df)
```


- Change orientation

```
sns.countplot(y = df['LABEL'], \n              hue=df['WEATHER1'], data = df)
```

- Save


```
myplot = sns.countplot(...
```

```
myplot.get_figure().savefig('plot.pdf')
```







Other common plot types



Box Plot
(1-dimensional)
`sns.boxplot`



Scatter Plot
(2-dimensional)
`sns.scatterplot`



Line Plot
(2-dimensional)
`sns.lineplot`